



Post-deployment tasks

HCI

NetApp

November 04, 2020

This PDF was generated from https://docs.netapp.com/us-en/hci/concept_nde_post_deployment_overview.html on November 04, 2020. Always check docs.netapp.com for the latest.

Table of Contents

- Post-deployment tasks 1
 - Post-deployment tasks 1
 - Supported networking changes 1
 - Disable the **smartd** service on NetApp HCI compute nodes 3
 - Keep VMware vSphere up to date 3
 - Install GPU drivers for GPU-enabled compute nodes 4
 - Configure Fully Qualified Domain Name web UI access 6
 - Access NetApp Hybrid Cloud Control 7

Post-deployment tasks

Post-deployment tasks

Depending on your choices during the deployment process, you need to complete some final tasks before your NetApp HCI system is ready for production use, such as updating firmware and drivers and making any needed final configuration changes.

Find more information

- [NetApp HCI Documentation Center](#)
- [SolidFire and Element Software Documentation Center](#)

Supported networking changes

After you deploy NetApp HCI, you can make limited changes to the default networking configuration. However, certain settings are required for smooth operation and proper network detection. Changing these settings will cause unexpected behavior, and might prevent you from expanding compute and storage resources.

After you deploy your system, you can make the following changes to the default network configuration in VMware vSphere as dictated by your network requirements:

- Change vSwitch names
- Change port group names
- Add and remove additional port groups
- Change the vmnic interface failover order for any additional port groups you have added

H300E, H500E, H700E and H410C compute nodes

NetApp HCI expects the following network configuration for H300E, H500E, H700E and H410C nodes.

The following is a six-interface configuration with VMware vSphere Distributed Switching (VDS). This configuration is only supported when used with VMware vSphere Distributed Switches, and requires VMware vSphere Enterprise Plus licensing.

| Network function | vmkernel | vmnic (physical interface) |
|------------------|----------|----------------------------------|
| Management | vmk0 | vmnic2 (Port A), vmnic3 (Port B) |

| Network function | vmkernel | vmnic (physical interface) |
|------------------|----------|----------------------------------|
| iSCSI-A | vmk1 | vmnic5 (Port E) |
| iSCSI-B | vmk2 | vmnic1 (Port D) |
| vMotion | vmk3 | vmnic4 (Port C), vmnic0 (Port F) |

The following is a six-interface configuration with VMware vSphere Standard Switching (VSS). This configuration uses VMware vSphere Standard Switches (VSS).

| Network function | vmkernel | vmnic (physical interface) |
|------------------|----------|----------------------------------|
| Management | vmk0 | vmnic1 (Port D), vmnic5 (Port E) |
| iSCSI-A | vmk1 | vmnic1 (Port E) |
| iSCSI-B | vmk2 | vmnic5 (Port D) |
| vMotion | vmk3 | vmnic1 (Port C), vmnic5 (Port F) |

H610C compute nodes

NetApp HCI expects the following network configuration for H610C nodes.

This configuration is only supported when used with VMware vSphere Distributed Switches (VDS), and requires VMware vSphere Enterprise Plus licensing.



Ports A and B are unused on the H610C.

| Network function | vmkernel | vmnic (physical interface) |
|------------------|----------|----------------------------------|
| Management | vmk0 | vmnic2 (Port C), vmnic3 (Port D) |
| iSCSI-A | vmk1 | vmnic3 (Port D) |
| iSCSI-B | vmk2 | vmnic2 (Port C) |
| vMotion | vmk3 | vmnic2 (Port C), vmnic3 (Port D) |

H615C compute nodes

NetApp HCI expects the following network configuration for H615C nodes.

This configuration is only supported when used with VMware vSphere Distributed Switches (VDS), and requires VMware vSphere Enterprise Plus licensing.

| Network function | vmkernel | vmnic (physical interface) |
|------------------|----------|----------------------------------|
| Management | vmk0 | vmnic0 (Port A), vmnic1 (Port B) |
| iSCSI-A | vmk1 | vmnic0 (Port B) |

| Network function | vmkernel | vmnic (physical interface) |
|------------------|----------|----------------------------------|
| iSCSI-B | vmk2 | vmnic1 (Port A) |
| vMotion | vmk3 | vmnic0 (Port A), vmnic1 (Port B) |

Find more information

- [NetApp HCI Documentation Center](#)
- [SolidFire and Element Software Documentation Center](#)

Disable the **smartd** service on NetApp HCI compute nodes

By default, the **smartd** service periodically polls the drives in your compute nodes. You should disable this service on all compute nodes after you deploy NetApp HCI.

Steps

1. Using SSH or a local console session, log in to VMware ESXi on the compute node using root credentials.
2. Stop the running **smartd** service:

```
/etc/init.d/smartd stop
```

3. Prevent the **smartd** service from starting at boot:

```
chkconfig smartd off
```

4. Repeat these steps on the rest of the compute nodes in your installation.

Find more information

- [Turn off the smartd service in VMware ESXi](#)
- [VMware KB article 2133286](#)

Keep VMware vSphere up to date

After deploying NetApp HCI, you should use VMware vSphere Lifecycle Manager to apply the latest security patches for the version of VMware vSphere used with NetApp HCI.

Use the [Interoperability Matrix Tool](#) to ensure that all versions of software are compatible. See the

[VMware vSphere Lifecycle Manager documentation](#) for more information.

Find more information

- [NetApp HCI Documentation Center](#)
- [SolidFire and Element Software Documentation Center](#)

Install GPU drivers for GPU-enabled compute nodes

Compute nodes with NVIDIA graphics processing units (GPUs), like the H610C, need NVIDIA software drivers installed in VMware ESXi so that they can take advantage of the increased processing power. After deploying compute nodes with GPUs, you need to perform these steps on each GPU-enabled compute node to install the GPU drivers in ESXi.

Steps

1. Open a browser and browse to the NVIDIA licensing portal at the following URL:

```
https://nvid.nvidia.com/dashboard/
```

2. Download one of the following driver packages to your computer, depending on your environment:

| vSphere version | Driver package |
|-----------------|--|
| vSphere 6.5 | NVIDIA-GRID-vSphere-6.5-410.92-410.91-412.16.zip |
| vSphere 6.7 | NVIDIA-GRID-vSphere-6.7-410.92-410.91-412.16.zip |

3. Extract the driver package on your computer.

The resulting .VIB file is the uncompressed driver file.

4. Copy the .VIB driver file from your computer to ESXi running on the compute node. The following example commands for each version assume that the driver is located in the `$HOME/NVIDIA/ESX6.x/` directory on the management host. The SCP utility is readily available in most Linux distributions, or available as a downloadable utility for all versions of Windows:

| ESXi version | Description |
|--------------|---|
| ESXi 6.5 | <pre>scp \$HOME/NVIDIA/ESX6.5/NVIDIA**.vib root@<ESXi_IP_ADDR>:/.</pre> |

| ESXi version | Description |
|--------------|--|
| ESXi 6.7 | <code>scp \$HOME/NVIDIA/ESX6.7/NVIDIA**.vib root@<ESXi_IP_ADDR>:/. </code> |

5. Use the following steps to log in as root to the ESXi host and install the NVIDIA vGPU Manager in ESXi.

a. Run the following command to log in to the ESXi host as the root user:

```
ssh root@<ESXi_IP_ADDRESS>
```

b. Run the following command to verify that no NVIDIA GPU drivers are currently installed:

```
nvidia-smi
```

This command should return the message `nvidia-smi: not found`.

c. Run the following commands to enable maintenance mode on the host and install the NVIDIA vGPU Manager from the VIB file:

```
esxcli system maintenanceMode set --enable true  
esxcli software vib install -v /NVIDIA**.vib
```

You should see the message `Operation finished successfully`.

d. Run the following command and verify that all eight GPU drivers are listed in the command output:

```
nvidia-smi
```

e. Run the following command to verify that the NVIDIA vGPU package was installed and loaded correctly:

```
vmkload_mod -l | grep nvidia
```

The command should return output similar to the following: `nvidia 816 13808`

f. Run the following command to reboot the host:

```
reboot -f
```

g. Run the following command to exit maintenance mode:

```
esxcli system maintenanceMode set --enable false
```

6. Repeat steps 4-6 for any other newly deployed compute nodes with NVIDIA GPUs.
7. Perform the following tasks using the instructions in the NVIDIA documentation site:
 - a. Install the NVIDIA license server.
 - b. Configure the virtual machine guests for NVIDIA vGPU software.
 - c. If you are using vGPU-enabled desktops in a virtual desktop infrastructure (VDI) context, configure VMware Horizon View for NVIDIA vGPU software.

Find more information

- [NetApp HCI Documentation Center](#)
- [SolidFire and Element Software Documentation Center](#)

Configure Fully Qualified Domain Name web UI access

NetApp HCI with Element 12.2 or later enables you to access storage cluster web interfaces using the Fully Qualified Domain Name (FQDN). If you want to use the FQDN to access web user interfaces such as the Element web UI, per-node UI, or management node UI, you must first add a storage cluster setting to identify the FQDN used by the cluster. This enables the cluster to properly redirect a login session and improves integration with external services such as key managers and identity providers for multi-factor authentication.

What you'll need

- This feature requires Element 12.2 or later and management services version 2.15 or later.

Steps

1. Create the following cluster interface preference using the `CreateClusterInterfacePreference` API method, inserting the cluster MVIP FQDN for the preference value:
 - Name: `mvip_fqdn`
 - Value: <Fully Qualified Domain Name for the Cluster MVIP>
2. Change the management node settings using the management node REST API:
 - a. Access the REST API UI for the management node by entering the management node IP address followed by `/mnode/2/`. For example: `https://[management node IP]/mnode/2/`
 - b. Click **Authorize** or any lock icon and complete the following:

- i. Enter the cluster user name and password.
- ii. Enter the client ID as `mnode-client`.
- iii. Click **Authorize** to begin a session.
- iv. Close the window.
- c. Click **GET /settings**.
- d. Click **Try it out**.
- e. Click **Execute**.
- f. Record any proxy settings reported in the response body.
- g. Click **PUT /settings**.
- h. Click **Try it out**.
- i. In the request body area, enter the management node FQDN as the value for the `mnode_fqdn` parameter.
- j. Enter any proxy setting values you recorded earlier in the remaining parameters in the request body. If you leave the proxy parameters empty or do not include them in the request body, existing proxy settings will be removed.
- k. Click **Execute**.

Access NetApp Hybrid Cloud Control

NetApp Hybrid Cloud Control enables you to manage NetApp HCI. You can upgrade management services and other components of NetApp HCI and expand and monitor your installation. You log in to NetApp Hybrid Cloud Control by browsing to the IP address of the management node.

What you'll need

You have upgraded your management services to at least version 2.1.326. NetApp Hybrid Cloud Control is not available in earlier service bundle versions. For information about the current service bundle version, see the [Management Services Release Notes](#).

Steps

1. Open a web browser and browse to the IP address of the management node. For example:

```
https://<ManagementNodeIP>
```

2. Log in to NetApp Hybrid Cloud Control by providing the NetApp HCI storage cluster administrator credentials.

The NetApp Hybrid Cloud Control interface appears.

Find more information

- [NetApp HCI Documentation Center](#)
- [SolidFire and Element Software Documentation Center](#)

Copyright Information

Copyright © 2020 NetApp, Inc. All rights reserved. Printed in the U.S. No part of this document covered by copyright may be reproduced in any form or by any means-graphic, electronic, or mechanical, including photocopying, recording, taping, or storage in an electronic retrieval system-without prior written permission of the copyright owner.

Software derived from copyrighted NetApp material is subject to the following license and disclaimer:

THIS SOFTWARE IS PROVIDED BY NETAPP “AS IS” AND WITHOUT ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE, WHICH ARE HEREBY DISCLAIMED. IN NO EVENT SHALL NETAPP BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

NetApp reserves the right to change any products described herein at any time, and without notice. NetApp assumes no responsibility or liability arising from the use of products described herein, except as expressly agreed to in writing by NetApp. The use or purchase of this product does not convey a license under any patent rights, trademark rights, or any other intellectual property rights of NetApp.

The product described in this manual may be protected by one or more U.S. patents, foreign patents, or pending applications.

RESTRICTED RIGHTS LEGEND: Use, duplication, or disclosure by the government is subject to restrictions as set forth in subparagraph (c)(1)(ii) of the Rights in Technical Data and Computer Software clause at DFARS 252.277-7103 (October 1988) and FAR 52-227-19 (June 1987).

Trademark Information

NETAPP, the NETAPP logo, and the marks listed at <http://www.netapp.com/TM> are trademarks of NetApp, Inc. Other company and product names may be trademarks of their respective owners.